Data Mining: Fundamentals

A.Y. 2024/2025

Group 12

Bruno Barbieri, Noemi Dalmasso, Gaia Federica Francesca Ferrara

Contents

1 Data Understanding and Preparation

The dataset *train.csv* contains 16431 titles of different forms of visual entertainment that have been rated on IMDb, an online database of information related to films, television series etc. Each record is described by 23 attributes, either discrete or continuous.

1.1 Discrete Attributes

Table 1.1 shows the discrete attributes of the dataset, their types and a brief description of each attribute.

Attribute	Type	Description
originalTitle	Categorical	Title in its original language
rating	Ordinal	IMDB title rating class, from (0,1] to (9,10]
worstRating	Ordinal	Worst title rating
bestRating	Ordinal	Best title rating
titleType	Categorical	The format of the title
canHaveEpisodes	Binary	Whether the title can have episodes: True/False
isRatable	Binary	Whether the title can be rated: True/False
isAdult	Binary	Whether the title is adult content: 0 (non-adult), 1 (adult)
countryOfOrigin	List	Countries where the title was produced
genres	List	Genre(s) associated with the title (up to 3)

Table 1.1: Description of discrete attributes

1.1.1 Merging and Removal of Discrete Attributes

The following discrete attributes were removed from the dataset:

- originalTitle was removed because it is not relevant for the analysis;

- the isRatable variable was removed because all the titles in the dataset are ratable;

- worstRating and bestRating attributes were removed because they assume the same values for all records (1 and 10 respectively).

Additionally, the isAdult attribute is highly correlated with the presence or absence of *Adult* in genre (16 records differ in the train set, 1 in the test set), so the two were merged with a logical OR operation. This is not true for the *short* type in titleType, with 491 records having different values from the obtained feature. For this reason, the two were kept separate.

1.1.2 Discrete Attributes Analysis

This paragraph provides an overview of the discrete attributes in the dataset, focusing on their distributions and statistics. Figures 1.1a and 1.1b show bar plots of titleType and rating attributes, respectively.
Figure 1.1a shows that the classes of the titleType attribute are imbalanced, with *movie* being the most frequent class (5535 records). It was observed that the class *tvShort* is the least frequent in the dataset, with only 40 records (around 0.24% of the dataset). Because of this, these rows were discarded from the dataset, as they were considered irrelevant for the analysis. The decision was not repeated for *tvSpecial* and *tvMiniSeries*, as they cover slightly more than 1% of the dataset each (166, 1.01% and 224, 1.36%, respectively).

(a) Distribution of `titleType`


(b) Distribution of `rating`

Figure 1.1: Distribution of the `titleType` and `rating` attributes

As shown in figure 1.1b, the `rating` attribute roughly follows a normal distribution, with a slightly asymmetric peak: a significant number of titles falls within the (6,7] and (7, 8] ranges (4565 and 4822 titles, respectively) while only a total amount of 67 titles falls within (0,1] and (1,2].

1.1.3 Encoding and Transformation of Categorical Attributes

The attribute `rating` was transformed by taking the upper bound of each rating interval's string representation. This approach was chosen because the minimum rating is 1, meaning the lowest interval corresponds only to ratings of 1. For consistency, the same transformation was applied to all other intervals. Multi-label one-hot encoding was applied to the `genres` column. Each unique genre was represented as a binary feature, allowing records that belong to multiple genres simultaneously to maintain this information; this generated 28 new features. Depending on the task, some were often discarded to avoid overfitting or to reduce the number of features. Rows with no genres were assigned a vector of all zeros, indicating the absence of any genres.

The attribute `countryOfOrigin` was represented by grouping the countries by continent. The following variables have been created:

- `countryOfOrigin_freq_enc`
  (frequency encoding of the original list);

- `countryOfOrigin_AF` (Africa);

- `countryOfOrigin_AS` (Asia);

- `countryOfOrigin_EU` (Europe);

- `countryOfOrigin_NA` (North America);

- `countryOfOrigin_SA` (South America);

- `countryOfOrigin_OC` (Oceania);

- `countryOfOrigin_UNK` (Unknown country).

For each record, the first six features provide the number of countries for each continent. The `countryOfOrigin_UNK` variable counts the number of countries that are not recognized as belonging to a continent for that record. Additionally, `countryOfOrigin_freq_enc` provides the frequency encoding of the original list of countries as a whole, showing how frequently a specific combination of countries appears across the entire dataset. These transformations allow to keep most of the original information, while limiting the number of new features.

1.2 Continuous Attributes

Table 1.2 shows the continuous attributes of the dataset, their type and a brief description.

Attribute	Type	Description
runtimeMinutes	Integer	Runtime of the title expressed in minutes
startYear	Integer	Release/start year of a title
endYear	Integer	TV Series end year
awardWins	Integer	Number of awards the title won
numVotes	Integer	Number of votes the title has received
totalImages	Integer	Number of images on the IMDb title page
totalVideos	Integer	Number of videos on the IMDb title page
totalCredits	Integer	Number of credits for the title
criticReviewsTotal	Integer	Total number of critic reviews
awardNominationsExcludeWins	Integer	Number of award nominations excluding wins
numRegions	Integer	Number of regions for this version of the title
userReviewsTotal	Integer	Number of user reviews
ratingCount	Integer	Total number of user ratings for the title

Table 1.2: Description of continuous attributes

1.2.1 Removal and Merging of Continuous Attributes

The plot in figure 1.2 is a Pearson's correlation matrix that takes into account the continuous attributes of the dataset. The matrix shows that `ratingCount` and `numVotes` are perfectly correlated; for their redundancy, `ratingCount` was discarded.


Figure 1.2: Correlation matrix

The attributes `awardNominationsExcludeWins` and `awardWins` were combined into `totalNominations`, due to their strong semantic similarity and high correlation (0.69). The new feature represents the sum of the two original attributes. This transformation also helps mitigate the impact of their heavy right skew (shown in figure1.3a), resulting in a more meaningful and interpretable feature. Similarly, the `totalVideos` and `totalImages` attributes were combined into a single feature, i.e. `totalMedia`, representing the total number of media items associated with a title. Although the original attributes are not highly correlated, both exhibit skewed distributions (as in figure1.3b), `totalVideos` in particular. Due to this, and to their similar semantic meaning, they were merged to form a more consolidated and interpretable feature.

(a) Kernel Density Estimation of `awardWins`
and `awardNominationsExcludeWins`

(b) Kernel Density Estimation of
`totalVideos` and `totalImages`

Figure 1.3: Distribution of the attributes that form the `totalNominations` and `totalMedia` features

Although `criticReviewsTotal` and `userReviewsTotal` also have a relatively high correlation (0.65), as well as a right-skewed distribution, it was decided that the two attributes should be kept separate because of their relevance in meaning. It is also worth noting that the two have high correlations with `numVotes` (0.67 and 0.75 respectively), but they were all kept because of the difference between votes and reviews.

1.3 Data Quality

Next, a proper evaluation of the observed data was conducted in preparation for the analysis. Once having checked that there are no duplicates and no incomplete rows in the dataset, attention was given at identifying missing values and outliers.

1.3.1 Missing Values

The following attributes were found to have missing values[1]:

- `endYear`: it is the feature with the highest number of `NaN` values (15617; about 95%). Although the feature is only relevant for *TVSeries* and *TVMiniSeries* titles, it still had approximately 50% missing values within those categories, limiting its usefulness even in the appropriate context. For this reason, the feature was discarded.

---

[1]missing values were marked as `NaN` only in `awardWins`; in the other listed columns "\N" was used, so it was then converted to `NaN`

- `runtimeMinutes`: this attribute has 4,852 missing values (29.5%). Two imputation strategies were employed, both based on random sampling within the interquartile range. The first used `titleType` to define the range, while the other imputed values based off of the attribute's distribution alone. The choice of which of the two strategies to use depends on the specific task, and will be specified in the corresponding sections.

- `awardWins`: this feature has 2618 `NaN` values (about 16%). Since the mode associated with this variable is 0, it has been decided to substitute the missing values with 0.

- `genres`: it has 382 missing values (2.3%). Having dealt this variable with a multi-label one-hot encoding process (as has been described in the *Encoding and Transformation of categorical attributes* section), a vector of all zeros is assigned to records with missing genres values.

1.3.2 Semantic Inconsistencies, Feature Transformations and Outlier detection

While analyzing the dataset, it was observed that the *Videogame* type of the `titleType` attribute (259 records - around 1.58% of the dataset) was not consistent with the other values of the same feature, being *Videogame* a fundamentally different type. These rows also generated problems for some of the other attributes, such as `runtimeMinutes`, resulting in most values being missing and difficult to impute. Because of this, the samples were removed from the dataset.

Some features showed a heavy right-skewed distribution, with typical traits of Power-Law Distributions. Their Kernel Density Estimations are shown in figure 1.4.

(a) KDE of right-skewed attributes   (b) KDE of right-skewed attributes with log density

Figure 1.4: Kernel Density Estimation of the right-skewed attributes

The decay of these features is exponential in linear space (1.4a), while in logarithmic space there is a decline that can be approximated to a linear trend (1.4b). For this reason, when needed, a log-transformation was applied to these attributes to reduce the skewness and make them more suitable for some specific analysis. Because of right-skewness (without a power-law distribution), other attributes were also log-transformed:

- `numVotes`;

- `totalCredits`;

- `totalMedia`.

Regarding outliers, `runtimeMinutes` was the most problematic feature. Similarly to missing values imputation (1.3.1), outlier detection was performed using two different strategies: the first approach computes outliers on each `titleType` separately, while the second is based on the distribution of the `runtimeMinutes` attribute alone. Figure 1.5 reports an analysis of the feature through the IQR method separately on each type. The boxplots show that there are samples that have been misreported, with runtimes of over 1000 minutes for *tvSeries*. This might be because of an inconsistency with the understanding of the meaning of the attribute, and in those cases it might be possible that the value refers to the total runtime of the series, rather than the runtime of a single episode. Another interesting observation regards the presence of a record with a runtime of 0 minutes for the *short* type; the record was removed because it was regarded as an erroneous sample. Other than this sample, other outliers were not removed from the dataset by default. Instead, a case-by-case approach was adopted, testing each task and analysis both with and without the outliers. Notably, in every case, better results were obtained when outliers were excluded.


Figure 1.5: Boxplot of the `runtimeMinutes` attribute for each `titleType`

2 Clustering

This chapter aims at illustrating the clustering analysis performed on the dataset at hand. The employed clustering techniques are K-means (Centroid-based), DBSCAN (density-based) and hierarchical clustering. The analysis is conducted using these methods on a selection of the continuous attributes of the dataset, which were appropriately log-transformed (when needed - as mentioned in subsection 1.3.2), and normalized using `MinMaxScaler`.

2.1 K-means

Clustering analysis with K-means was performed using a carefully selected subset of features. This step was motivated by the algorithm's sensitivity to the curse of dimensionality, as including too many variables can negatively affect SSE and Silhouette scores. For this reason, although only `numVotes`, `totalCredits`, `userReviewsTotal`, `runtimeMinutes`, and `criticReviewsTotal`, were chosen for this task, they proved to be a solid choice due to their ability to represent meaningful aspects of the data. To identify the proper number of clusters, both SSE and Silhouette scores were evaluated. The objective was to find a configuration that reduces the SSE while maintaining a robust Silhouette score and a proper k.

The plots in Figure 2.1a show that $k = 4$ provides a balance between the two values (obtaining SSE equal to 535.51 and 0.315 for the Silhouette score). Only for visualization purposes, Principal Component Analysis (PCA) was employed. The first two components account for 58.01% and 21.48% of the total variance, indicating that this projection provides a fairly informative view of the clustering structure. The resulting clusters are presented in figure 2.1b. The 4 distinct clusters appear well-separated, suggesting that K-means managed to capture meaningful groupings in the data.
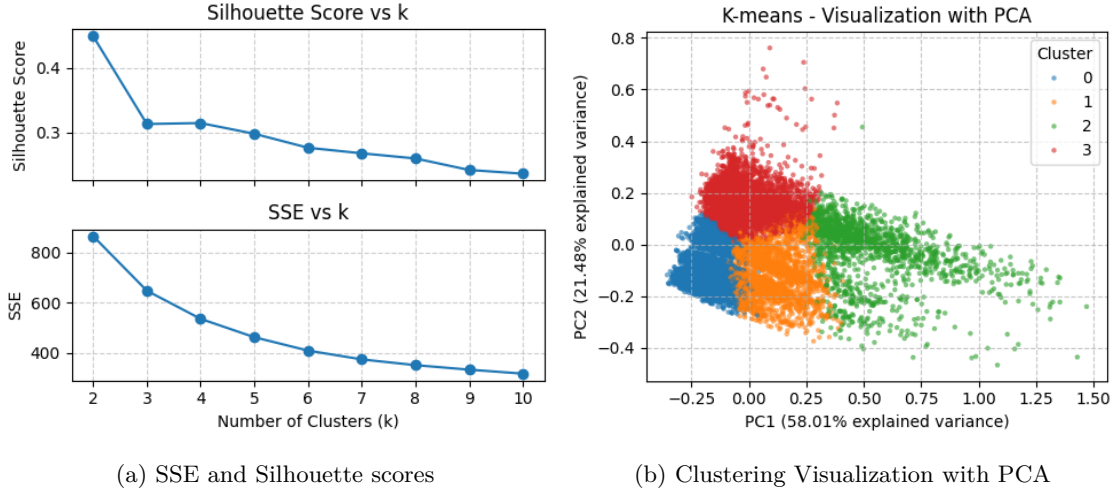
(a) SSE and Silhouette scores

(b) Clustering Visualization with PCA

Figure 2.1: K-means - Cluster analysis

2.2 DBSCAN

To determine suitable DBSCAN parameters, the k-th distance plot was used as starting point (Figure 2.2a). By varying k (*MinPts*) and observing the corresponding distance curves, the range between 0.08 and 0.2 was identified for a possible *Eps* value. The algorithm was applied to the following features: `numVotes`, `totalCredits`, `criticReviewsTotal`, `userReviewsTotal`, and `runtimeMinutes`. Although different combinations were tested, they all led to similar results. Effectively, initial experiments clearly revealed DBSCAN's sensitivity to highly sparse and imbalanced data; in most cases, the algorithm produced a highly populated single cluster corresponding to the dominant dense region, while leaving the sparser regions largely unstructured.

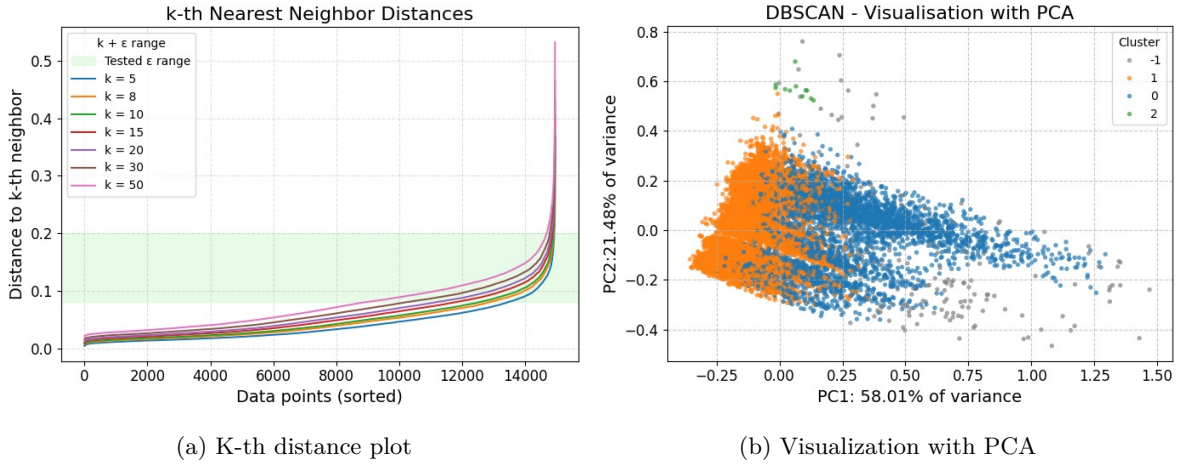(a) K-th distance plot

(b) Visualization with PCA

Figure 2.2: DBSCAN - Cluster analysis

To address this issue and obtain more meaningful clustering results, different combinations of `min_eps` and `min_points` in the following ranges were tested:

- `eps_list` = [0.08, 0.09, 0.1, 0.12, 0.14, 0.16, 0.18, 0.2, 0.3, 0.4, 1.0];

- `min_points_list` = [4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 20].

Each configuration was evaluated based on the number of clusters (with a minimum threshold of 2), their size, noise points, and Silhouette score (filtered for values > 0.1), excluding cases with a single cluster or very low Silhouette values. Although results were not fully satisfactory, the most relevant configurations are summarized in Table 2.1. The reported Silhouette scores are computed excluding noise points; therefore, no score is given when only one cluster is formed. Note that, in this context, high Silhouette scores can be misleading, as they can favor compactness over structural granularity,

eps threshold	Min points	Silhouette	Cluster composition
$\varepsilon > 0.16$	for any tested value	–	never led to more than one cluster
$0.13 \leq \varepsilon \leq 0.16$	for any tested value	[0.54, 0.70]	led to one cluster and, occasionally, 1 or 2 minor ones
$\varepsilon < 0.13$	for any tested value	[0.10, 0.31]	the big cluster it's splitted in two different ones
$\varepsilon \leq 0.10$	lower values [4, 5]	[0.10, 0.25]	led to several, not meaningful, clusters of 5/6 points

Table 2.1: Relevant DBSCAN parameters configurations

Excluding the combination with higher *Eps* value (that made the points collapse into one cluster), and the ones with low *Eps* and few `min_pts` (that resulted in over-fragmentation) the analysis focused on *Eps* in the range [0.11, 0.13]. Despite the limitations of the dataset structure, in this range the algorithm was able to differentiate between regions of varying density.

One of the most interesting compromise (though still suboptimal in absolute terms) was found with *Eps* = 0.11, `min_points` = 7: As shown in (Figure 2.2b), this setting resulted in a three-cluster structure. A compact one of 10378 points was extracted from the densest region. A second less dense cluster with fewer points (4347) captured intermediate-density areas or mild outliers. Finally, a very small but well-separated cluster, of only 11 points, effectively captured a potential relevant local pattern. In addition, 219 points were labeled as noise. The resulting Silhouette Score was 0.2699.

2.3 Hierarchical clustering

Hierarchical clustering was performed using all linkages (Ward, Average, Complete, Single), with the Euclidean distance metric. After a careful analysis of multiple clusterings, it was decided to procede with all remaining numerical features. Figure 2.3 shows the results of the analysis, which includes the Silhouette and SSE scores, as well as the maximum and minimum percentage of points per cluster.

From figure 2.3a, it can be observed that Single linkage produces a single cluster which contains basically all data points. This makes it unsuitable for this use case. Average and Complete linkages produce a cluster with a high maximum cluster size (above 90% of the dataset for most of the number of clusters tested for Average, and 80% for Complete). These results are likely due to two main causes:

- Usage of skewed features, which lead to areas with a very high density of points;

- High dimensionality of the dataset, which makes it more difficult to separate clusters effectively.

These issues are mitigated by Ward's method, which doesn't show a dominant biggest cluster, for all numbers of clusters tested. The minimum cluster size is also consistently more balanced across different numbers of clusters.

(a) Max/Min percentage of points per cluster   (b) Silhouette and SSE scores
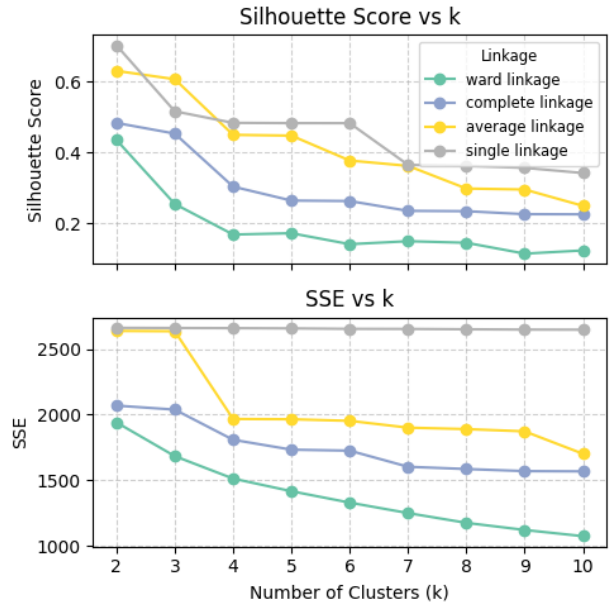
Figure 2.3: Hierarchical clustering metrics for different numbers of clusters

Since Ward's method is based on Squared Error, it's not surprising how its SSE is consistently lower than other linkages, as shown in figure 2.3b. More interesting is that Ward's method has the lowest Silhouette scores for all numbers of clusters tested, which indicates that the clusters are not well separated. This is due to the fact that Ward's method groups the data which resides in the high-density area mentioned above into smaller clusters, leading to less well-defined boundaries between them. In the next two sections, the results of Ward's method and Complete Linkage will be analyzed; the other two will not be discussed further, as they provide limited insights on the dataset.

2.3.1 Ward's method

(a) Ward's Method Dendrogram   (b) Ward's Method Scatter plot

Figure 2.4: Ward's method clustering

9

Figure 2.4 shows a dendrogram and a scatter plot of the clustering obtained through Ward's Method. The cut at 4 clusters was chosen based on the parameters discussed in the previous section and on the clarity of the dendrogram.

The dendrogram in figure 2.4a reveals well-separated clusters, all of which merge at a similar linkage distance (approximately 12). This indicates that the increase in within-cluster sum of squares (SSE) is relatively consistent across all of the final clusters' merges. Notably, the smallest cluster is the last to be merged, which suggests it is more distinct from the others. As shown in Figure 2.4b, this cluster lies in a sparser region of the dataset. In contrast, the remaining three clusters originate from the denser region, and are therefore spatially closer to one another. Compared to K-Means, the clusters identified by Ward's method appear less clearly separated in the PCA projection, resulting in some overlap between adjacent groups.

2.3.2   Complete Linkage

Figure 2.5 shows the dendrogram and scatter plot for a clustering obtained through Complete Linkage. Since the tendency of this linkage is to merge all points into a single cluster, the clustering is cut at 4 clusters, which helps mitigating the issue. While selecting five clusters would have introduced an additional split within the largest cluster with similar SSE and Silhouette scores, the resulting group was found to be poorly separated and lacked meaningful distinction. As such, it was not considered a valuable contribution to the overall clustering structure.

As shown in the dendrogram in Figure 2.5a, the clusters merge at similar linkage distances (approximately between 1.2 and 1.4), indicating that the maximum within-cluster distances are comparable across clusters. With respect to the clustering obtained through Ward's method, the clusters have clearer boundaries along the PCA axes, as the denser area of the dataset is not split into multiple clusters.

It is also interesting to observe how the dendrogram structure differs from that of Ward's method. In the previous case, the smallest cluster was the last to be merged, reflecting its distinctiveness in terms of within-cluster variance. In contrast, the Complete Linkage dendrogram shows the two smaller clusters being merged before the root. This is observed with Single and Average Linkages as well, and is a product of the sparsity of these clusters' regions.
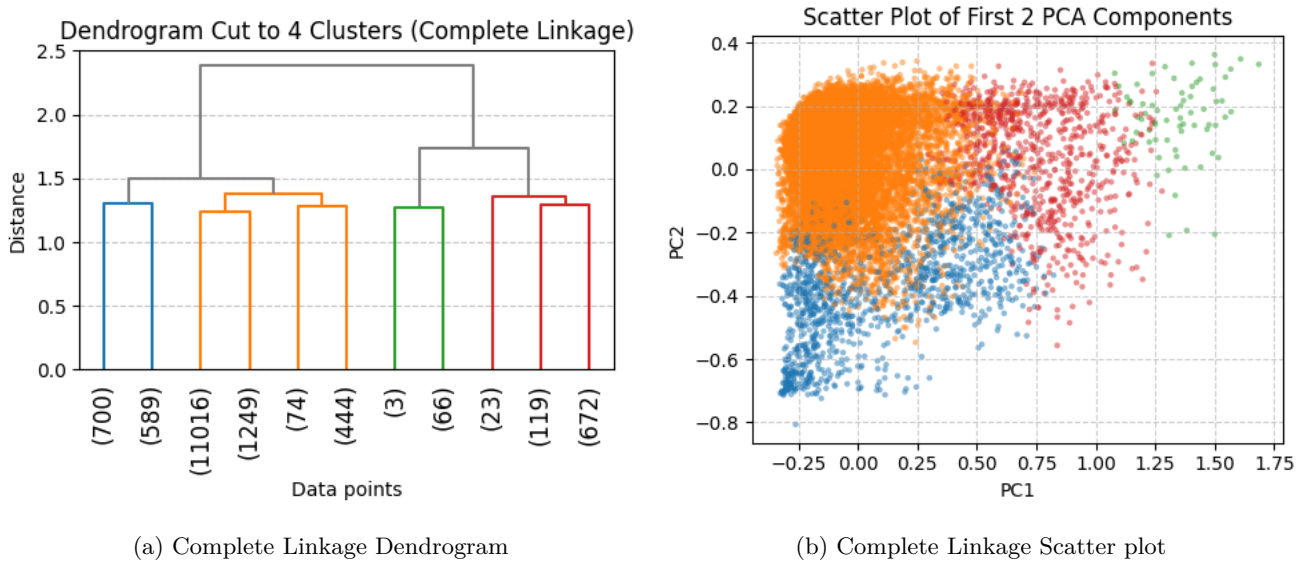

(a) Complete Linkage Dendrogram   (b) Complete Linkage Scatter plot

Figure 2.5: Dendrograms for hierarchical clustering with Complete Linkage

2.4 General considerations

In terms of performance metrics, K-means achieved the highest Silhouette score (0.315) and lowest SSE (535.51). This was possible after reducing the number of variables. This reflects the algorithm's tendency to perform well in simplified feature spaces, though at the cost of losing some informational richness.

In comparison, hierarchical clustering using Ward's method, which also minimizes SSE, led to worse performance metrics (Silhouette 0.17, SSE 1545) and poor separation between clusters, but was able to retain more information by using more variables. Complete Linkage had a similar SSE (1800), but a much higher Silhoette score (0.30), resulting in better separated clusters. However, it led to a heavily imbalanced cluster distribution, with one cluster containing over 80% of the data points.

DBSCAN produced acceptable Silhouette scores as well (0.269, excluding noise), but this is largely due to its tendency to merge dense areas together, while treating isolated points as noise. While this can improve internal cohesion metrics, it does not reflect meaningful segmentation, especially when the data contains a densely dominant area. However, the analysis revealed parameter ranges where meaningful subclusters start to emerge, beyond the dominant density mass.

Overall, each algorithm showed trade-offs: K-means achieved compact clusters with fewer variables, hierarchical clustering retained more information but resulted in eitherless separation or imbalanced clusters, and DBSCAN captured dense areas but was sensitive to data structure and tuning.

3 Classification

Classification was performed on the available training set using three different algorithms: K-NN (*K-Nearest Neighbours*), Naïve Bayes and Decision Tree. A binary classification task and a multiclass one were performed, respectively using the target variables `has_LowEngagement` and `titleType`. These were chosen because of their importance in terms of semantic meaning, and centrality in the dataset's structure.

3.1 Binary classification

The binary target variable used in this task, `has_LowEngagement`, was specifically defined for this purpose. It identifies records where the `numVotes` attribute is less than 100.

An analysis of semantically related features was run to decide whether to discard any. `userReviewsTotal` has a 75% correlation with `numVotes`, while `criticReviewsTotal` has 67% correlation. Due to its semantic similarity with `numVotes`, `userReviewsTotal` was discarded for this task. In contrast, `criticReviewsTotal` was retained as it was deemed to provide distinct and complementary information. Moreover, its correlation with other features was not considered sufficiently high to make the problem trivial. An important aspect of the chosen binary classification task is the class imbalance, with 10287 records in the training set belonging to the *Low Engagement* class and 4668 to *High Engagement*. This imbalance was taken into account during model training and evaluation, with a focus on macro-averaged F1-score to mitigate its impact on the results.

3.1.1 K-NN - Binary Classification

The features used for the K-NN algorithm were normalized, as the model is sensitive to unscaled values: log-transformation (when needed) and `StandardScaler` were applied to data. Then, to perform this task, care was taken in selecting the appropriate type and amount of features to avoid unnecessary increased dimensionality.

For example, only `countryOfOrigin_freq_enc` was used instead of including all 7 `countryOfOrigin_[continent code]` attributes. The remaining variables of the dataset were maintained as they offered diverse and non-redundant insights.

For this algorithm, a portion of the training set was held out as a validation set to enable external validation after internal cross-validation. A randomized hyperparameter search was indeed performed on the reduced training set, with each configuration evaluated using stratified 5-fold cross-validation. The objective was to identify the optimal value of k and explore additional algorithm parameters. The configuration that achieved the highest accuracy on the selected feature set is: `weights:'uniform', n_neighbors:9, metric:'cityblock'`.
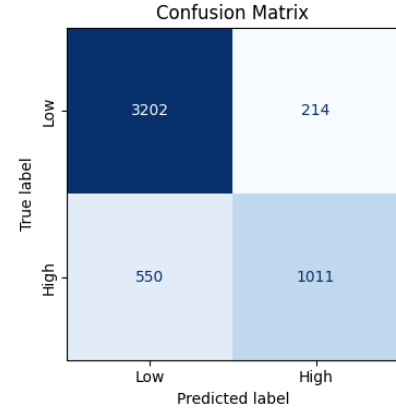

Figure 3.1: K-NN binary classification

The resulting model shows solid performance, achieving a test accuracy of 0.85. As expected, it performs better on the *Low Engagement* class, with a high recall (0.94) and F1-score (0.89). In contrast, the *High Engagement* class is more challenging for the model; while precision is reasonable (0.83), recall (0.65) indicates that a significant number of instances are misclassified as *Low Engagement*. This analysis is reflected in the confusion matrix in Figure 3.1, where 550 out of 1561 *High Engagement* samples were incorrectly labeled as *Low Engagement*. In general, the macro average F1-score of 0.81 suggests that the model maintains a balanced performance across both classes.

### 3.1.2 Naïve Bayes - Binary Classification

Figure 3.2: NB binary classification

For Naïve Bayes, a validation set was used to select the best feature set. Two variants of the algorithm were evaluated: GaussianNB and CategoricalNB. However, skewed distribution of continuous attributes and the relevance of discrete features (e.g., genres), made GaussianNB less suitable for the task, resulting in a slightly lower performance in terms of accuracy and classification of underrepresented classes. Therefore, only CategoricalNB results are reported. To apply CategoricalNB, continuous features were discretized using meaningful bins, and each `countryOfOrigin_[continent code]` feature was binarized (0/>0) to maintain only the significant information; however, only `is_from_America_bin` and `is_from_Europe_bin` were used. All features were encoded with OrdinalEncoder as required. The model achieved an overall accuracy of 0.83: a solid outcome for a binary classification task considering the simplicity of the algorithm.

However, differences between classes emerged: the *Low Engagement* class performed better, with a precision of 0.87, a recall of 0.89, and an F1-score of 0.88. On the other hand, the model struggled more in identifying *High Engagement* examples, as confirmed by the low recall for that class (0.70) and by the 481 false negatives in the confusion matrix (Figure 3.2). The macro average F1-score of 0.80 actually indicates a good balance overall.

3.1.3 Decision Tree - Binary Classification

For explainability purposes, features were not normalized nor transformed for the Decision Tree model, as it does not require such preprocessing because it's not based on distance measures, but rather on decision thresholds.

To identify the optimal hyperparameters, a Randomized Search was performed using Repeated Stratified 5-Fold Cross-Validation with 10 repeats on the training set, optimized for the macro-averaged F1-score. The best configuration found used Gini index as the splitting criterion, a maximum tree depth of 26, and a minimum of 3 samples per leaf. The obtained decision tree is shown in Figure 3.3.
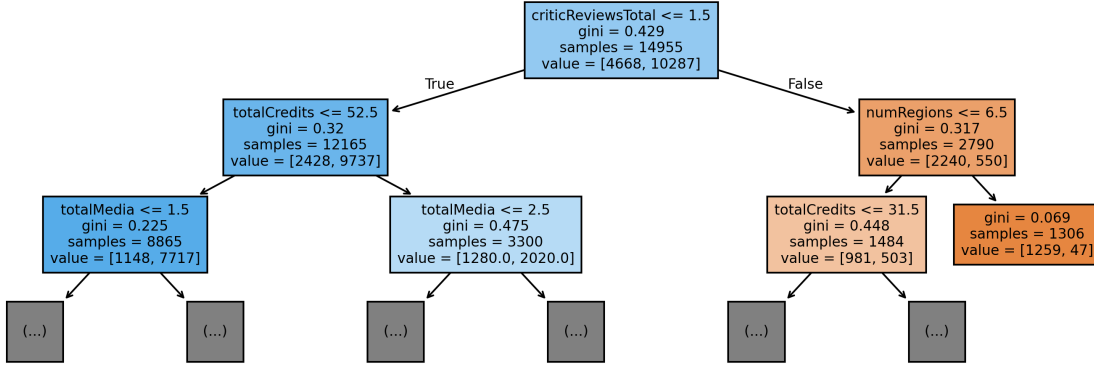

Figure 3.3: Decision Tree for binary classification

Unsurprisingly, the most important feature for the model was `criticReviewsTotal`, which amounted to 0.6 in the feature importance ranking. The 3 following more important features were `totalCredits` (0.15), `totalMedia` (0.10) and `numRegions` (0.09). These four take up around 93% of the total feature importance, and are all present in the first two splits shown in the Decision Tree.

Train performance was overall similar to the test one; in particular, the respective accuracies were of 0.83 and 0.84, and macro-F1 scores were 0.82 and 0.80. In any case, post-pruning was tested, but did not yield any performance improvement, and was therefore not applied. The *High Engagement* class showed low recall values (0.71 on train set, 0.68 on test set). This might be a consequence of class imbalance, as well as poor separability of the two classes. This assumption is further supported by the precision scores of the class (0.78 on train, 0.75 on test).

3.1.4 Model Comparison - Binary Classification

After applying the three algorithms to the binary classification task, it emerged that all models struggle to identify the minority class, *High Engagement*, leading to lower recall values: 0.65 for K-NN, 0.70 for Naïve Bayes, and 0.69 for Decision Tree. This is confirmed by the confusion matrices, where it appears that many *High Engagement* records are misclassified as *Low Engagement*.

Among the models, K-NN handled the minority class best in terms of precision (0.83) and F1-score (0.73). However, the fact that Naïve Bayes performed the highest in terms of recall (0.70) makes it a more valid alternative in scenarios where preventing high-engagement titles from being misclassified as low-engagement ones is a priority. In general, Decision Tree show similar performances to Naïve Bayes for all metrics In conclusion, K-NN is still considered the most efficient model for this task, as it also achieves the highest overall accuracy, the best macro-average F1-score and the most balanced results across both classes. The macro-average ROC AUC scores for all three models are also quite similar, with K-NN showing the highest value (0.89), followed by Naïve Bayes (0.88) and Decision Tree (0.87). A summary of the results of the best model is presented in Table 3.1.

	Accuracy	Precision Macro Avg	Recall Macro Avg	F1-score Macro Avg	AUC
K-NN	0.85	0.84	0.79	0.81	0.89

Table 3.1: Overall metrics - K-NN

3.2 Multiclass classification

Among the multiclass features in the training set, `titleType` was selected as the target variable for this task, due to its relevance within the dataset. Because of their strong correlation with `titleType`, the feature `canHaveEpisodes` and the genre *Short* were excluded from the feature set. Furthermore, since the primary imputation method for missing values in `runtimeMinutes` relied on information from the chosen target variable, these values were re-imputed to avoid data leakage. Specifically, missing entries were filled by sampling from the overall distribution of `runtimeMinutes`, without referencing `titleType` itself. For the same reason, outliers detected by observing the distribution of `runtimeMinutes` were removed, and not the ones based on `titleType` (see subsection 1.3.2).

One final point to note is the imbalance in the target feature (previously shown in figure 1.1a), which was explicitly taken into account during the design of the models. As for the binary classification task, macro-averaged F1-score was a key metric for model evaluation, as it provides a balance between each class's precision and recall.

3.2.1 K-NN - Multiclass Classification

Figure 3.4: K-NN confusion matrix

For the multiclass classification task using the K-NN algorithm, the inclusion of a broad set of features is justified, as it helps the model to better distinguish between underrepresented classes, despite not reaching optimal results. The hyperparameter tuning strategy that was employed is the same as in binary classification with K-NN. The configuration that achieved the best performances is: `weights:'uniform', n_neighbors:12, metric:'cityblock'`.

As expected, the model shows mixed performances across the different classes. The overall test accuracy reaches 0.76, which is a respectable result for a multiclass problem. However, since accuracy is a global metric, more attention was paid to the performances of the single classes. The ones with larger support, like *tvEpisode* (4690 records in the training set) and *movie* (5442), are handled quite well: *tvEpisode* reaches an F1-score of 0.83 and *movie* 0.84, both supported by high recall values (respectively 0.93 and 0.87).

This is also clearly visible in the confusion matrix of Figure 3.4, where these classes dominate with a high count of correct predictions. On the other hand, the model struggles with 2 classes in particular: *tvMiniSeries* and *tvSpecial*, i.e. the ones with the least support (respectively 186 and 158). Interesingly, the matrix shows how records belonging to class *tvSeries* are frequently misclassified as *tvEpisodes*; this is likely due to the similar nature of the two. This is reflected in a very low recall (0.24) for *tvSeries* and in a quite high precision (0.76) for *tvEpisode*.

To summarize, the macro average F1-score of 0.50 reflects clearly the imbalance in performance: while the model performs well on the majority classes, it fails to generalize well especially the minority ones.

3.2.2 Naïve Bayes - Multiclass Classification

For the multiclass classification task, the same procedures as binary classification (explained in subsection 3.1.2) were applied. Note that including the one-hot encoded *genres* features significantly improved classification across all classes, particularly those with limited support in the training set. While with GaussianNB, *tvMiniSeries* was not classified at all, and *tvMovie* and *tvSpecial* were poorly classified. On the other hand, since CategoricalNB better handles categorical data, these classes showed improvements, confirming its suitability for this task.

The model achieves an overall accuracy of 0.73, reasonable for multiclass classification, but indicative of difficulties in handling all categories equally. For frequent classes like *movie*, *tvEpisode*, and *short* high F1-scores (about 0.80) suggest effective recognition, likely due also to their distinctive traits. In addition, the *movie* class also shows strong precision (0.82) and recall (0.84), confirming the model's reliability on this category. As expected, also for NB, performance drops for underrepresented classes: *tvMiniSeries*, *tvSpecial*, and *tvMovie* have low F1-scores (respectively 0.02, 0.21, and 0.17). Other than their low support, this is probably linked to the presence of common features with more dominant classes, such as *tvEpisode* or *movie* (e.g., 140 *tvMovie* instances misclassified as movie). Also the *tvSeries* class shows issues, with moderate precision (0.52) but low recall (0.40), indicating many missed instances. The very low macro F1-score (0.46) actually reflects the inability of the model to perform in a balanced way on all the classes.

Figure 3.5: NB confusion matrix

3.2.3 Decision Tree - Multiclass Classification

Regarding feature and hyperparameter selection, the approach for this task follows the same setup used in the binary classification. The best configuration found used Entropy as splitting criterion, had a max depth of 14, a minimum of 5 samples per leaf, and a minimum of 8 samples in order to split an internal node.
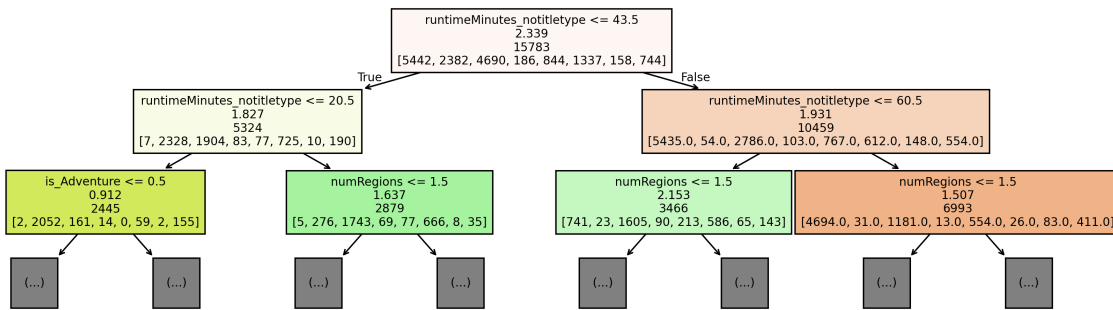
Figure 3.6: Decision Tree for multiclass classification

Figure 3.6 shows the Decision Tree obtained for the multiclass classification task. The most important feature for the model was `runtimeMinutes`, on which the first two levels of the tree are based, with a feature importance of 0.45. Three out of the four splits in the following level are based on `numRegions` being less than 2, giving the feature an importance of 0.13; the fourth is based on the *Adventure* genre, which has an importance of 0.01.

The model showed a general tendency towards overfitting; many configurations were tested to prevent this. The overall accuracy shows a significant drop (from 0.86 on the train set, to 0.75 on the test set), as well as the macro-averaged F1-score, going from 0.65 to 0.52. This was given from the *tvMiniSeries* and *tvSpecial* classes: on the train set these had f1-scores of respectively 0.35 and 0.22, and both had 0.10 on the test set. This can be seen in Figure 3.7, which shows that most predictions for these classes were wrong.

Post-pruning was also tested, but did not give particular benefits. Higher values for the parameter α had minor positive effects on the generalization capabilities of the model, at the cost of losing predictions on the less represented classes.

Figure 3.7: DT Confusion matrix


The confusion matrix also highlights that *tvMovie* was often misclassified as *movie*, resulting in a recall of 0.26 on the test set. This is likely due to semantic overlap between the two classes, as it is supported by the fact that the most frequent misclassification for *movie* was *tvMovie* itself, though it occurred less frequently, being movie is the largest class. A similar pattern is observed with *tvSeries* and *tvMiniSeries*, with 37 out of the total 68 of the *tvMiniSeries* records being misclassified as *tvSeries*.

3.2.4 Model Comparison - Multiclass Classification

Figure 3.8: Model comparison - ROC

The analysis of K-NN, CategoricalNB and Decision Tree models reveals several consistent patterns in their performances. As expected, underrepresented classes show consistently low recall and F1-scores across all models. *tvMiniSeries* performs worst, with recall between 0.01 (Naïve Bayes) and 0.07 (Decision Tree), and a maximum F1-score of 0.10. In contrast, well-represented classes like *movie* (5442) and *tvEpisode* (4690) are always classified accurately.

However, class imbalance alone does not fully explain the misclassifications. The confusion matrices showed systematic misclassification patterns, where minority classes were frequently absorbed into semantically related but more dominant categories. Across all models, but at different extent, *tvMovie* is frequently confused with *movie*, and *tvMiniSeries* is often predicted as *tvSeries*. K-NN also shows a tendency to predict *tvSeries* as *tvEpisode*. The choice of keeping those underrepresented classes, even if it lowers the overall performance (accuracy) of the models, is an important trade-off in favor of larger coverage and sensitivity across all classes.

To summarize, Decision Tree achieved the best overall performance (0.79 accuracy, 0.52 macro-average F1), followed by K-NN and Naïve Bayes. However, the ROC analysis reveals an apparent paradox: Naïve Bayes showed the highest macro-average AUC (0.91), despite poorer classification results. Still, Decision Tree remains the optimal choice, effectively balancing overall accuracy with minority class performance, since relying only on AUC can be misleading in imbalanced multiclass scenarios. In Table 3.2 its overall results are summarized.

	Accuracy	Precision Macro Avg	Recall Macro Avg	F1-score Macro Avg	AUC
DT	0.79	0.55	0.51	0.52	0.84

Table 3.2: Overall metrics - Decision Tree

4 Regression

Different regression techniques were applied to the dataset, testing on different combinations of attributes. The target variable chosen for Univariate and Multiple regression was `criticReviewsTotal`, while for Multivariate regression, the target variables were both `userReviewsTotal` and `criticReviewsTotal`. These were chosen because they offer important insights into the engagement that a product can generate, which also was the focus of the binary classification task in section 3.1.

4.1 Univariate and Multiple Regression

For Univariate Regression, the attribute `criticReviewsTotal` was chosen as the target variable. Aside from the semantic meaning, this choice was also made because it has a high correlation with the attribute `userReviewsTotal`, allowing univariate regression to be performed, while maintaining a clear separate semantic meaning. Multivariate Regression was performed with `numVotes`, `numRegions` and `totalMedia` as additional features, because of their high correlations with the target variable. All models' parameters were optimized using Cross-Validation, using Negative Mean Squared Error as scoring criterion. Table 4.1 shows the test performances of the different regressors.

	R^2	MAE	MSE
Univariate			
Linear	0.465	3.419	189.543
Ridge	0.465	3.419	189.550
Lasso	0.279	3.948	255.732
DT	0.700	2.313	106.531
47-NN	0.643	2.221	126.539
Multiple			
Linear	0.628	3.125	132.002
Ridge	0.627	3.071	132.119
Lasso	0.610	2.731	138.249
DT	0.702	2.119	105.854
15-NN	0.673	2.268	105.864

Table 4.1: Performance report for Univariate and Multiple Regression

The linear models underperformed compared to both Decision Tree and K-Nearest Neighbors regressors in both tasks. In the first regression, neither L1 nor L2 regularization improved performances, with Lasso performing worse than the other linear models. The same goes for Multiple Regression, but in this case L1 regularization had a slightly lower Mean Absolute Error than the other linear models. All linear models had a significant boost in all performance metrics with the additional features in Multiple Regression. Here, although performances were still worse than the other two, Lasso regression had a lower Mean Absolute Error.

K-NN regressors had solid performances, with a slight increase in Multiple Regression for the Mean Absolute Error, but a significantly lower Mean Squared Error. The best performing model in both tasks was the Decision Tree regressor, although the additional features offered a slight improvement in performance. An interesting observation is that the Cross-Validation best hyperparameters for Univariate Regression indicated a high maximum depth of 35 for Univariate Regression, while the best hyperparameter for Multiple Regression was a maximum depth of 8; a similar trend was observed in the K-NN regressors, where the optimal number of neighbors was 47 for Univariate Regression and 15 for Multiple Regression. This suggests that the additional features in Multiple Regression allowed the model to generalize better, reducing the need for a more complex model.

4.2   Multivariate Regression

As for Univariate and Multiple regressions, independent variables were chosen by studying the correlations of other features with the target variables. The chosen target variables (`userReviewsTotal` and `criticReviewsTotal`) shared the same top correlated features, which were found to be `totalCredits`, `numRegions` and `totalMedia`. As in previous regression tasks, hyperparameters were optimized using Cross-Validation, using Negative Mean Squared Error as scoring criterion . Table 4.2 shows the test performances of the different regressors.

	user R^2	user MAE	user MSE	critic R^2	critic MAE	critic MAE
Linear	0.812	5.877	721.100	0.623	3.339	131.068
Ridge	0.812	5.877	721.099	0.623	3.339	131.069
Lasso	0.812	5.624	722.818	0.604	3.113	137.847
DT	0.781	4.100	841.425	0.706	2.247	102.348
22-NN	0.761	4.268	917.439	0.685	2.278	109.629

Table 4.2: Performance report for Multivariate Regression

The linear models had generally similar performances. Regularization seemed to provide some improvements on some metrics, but overall the differences were not particularly significant, and both α parameters were kept low by the Cross-Validation. This suggests that the Linear model did not need particular regularization to perform at its best. Similarly to the previous task, the linear models had worse performances than the Decision Tree and K-NN regressors on the `criticReviewsTotal` target variable. The performance on this variable shows that all models achieved lower R^2 scores compared to those for `userReviewsTotal`, suggesting that its variance is more difficult to explain.

For `userReviewsTotal`, all linear models achieved higher R^2 scores than the Decision Tree and K-Nearest Neighbors models, with the Lasso and Linear regressors performing best and yielding lower Mean Squared Errors. However, their Mean Absolute Errors were higher than those of both the Decision Tree and K-NN regressors. This suggests that the linear models were able to capture a substantial portion of the variance in the target variable, but struggled to provide accurate predictions for individual titles. This may be due to the presence of outliers and noise in the dataset, as indicated by the generally high Mean Squared Error across all models. Additionally, the linear models may be limited by underlying non-linear relationships between features and the target variable. This last hypothesis is further supported by the high complexity of the Decision Tree (maximum depth of 45) and the K-NN regressor (which considered 22 neighbors), both of which adapted to more complex patterns in the data.

5   Pattern Mining

To perform this task, continuous attributes were discretized based on their distributions, aiming for bins that were both semantically meaningful and reasonably balanced in size. The selected numerical attributes (not normalized) for the pattern mining task, along with their binning, are:

Attribute	Binning
runtimeMinutes	VeryLowRTM (1-30), LowRTM (31-60), MediumRTM (61-90), HighRTM (91-220)
rating	VeryLowR (1-3), LowR (4-6), MediumR (7), HighR (8), VeryHighR (9-10)
totalCredits	VeryLowC (0-15), LowC (16-35), MediumC (36-65), HighC (66-15742)

Table 5.1: Binning of the continuous attributes

Regarding discrete attributes, `titleType` was considered for this task, kept in its original form. On the other hand, the values of each of the 7 attributes `countryOfOrigin_[continent code]`[1] were binarized into:

- `not_from_[continent code]`: when the value of the attribute is 0
- `is_from_[continent code]`: when the value of the attribute is ≥ 1

Binarization was preferred over creation of multiple bins, as it allows for a more linear interpretation of the results. An attempt was performed to include genres; however, this did not lead to more interesting results.

5.1   Extraction and discussion of frequent patterns

(a) Plot of frequent itemsets with varying *support*

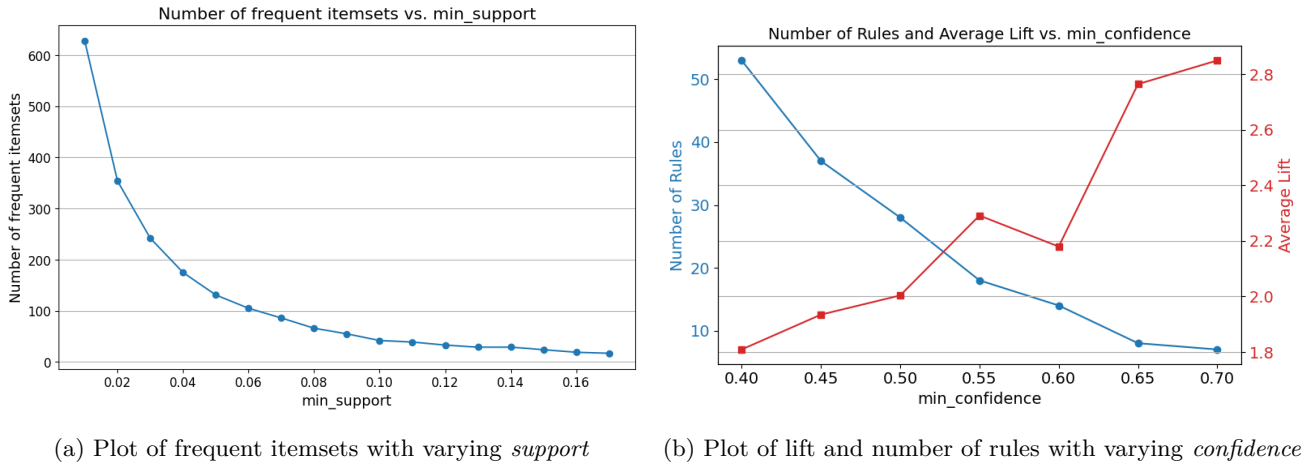(b) Plot of lift and number of rules with varying *confidence*

Figure 5.1: Plots of minimum support and confidence - *Apriori* algorithm

Apriori was the algorithm chosen to extract frequent patterns. Figure 5.1a shows how the number of frequent itemsets changes with support values ranging from 0.01 to 0.18. The curve of the plot begins to flatten between 0.08 and 0.1, so a support value of 0.08 was selected, resulting in 66 frequent patterns.

It is interesting to observe the top frequent itemsets of size 1, 2, and 3, as shown in Table 5.2. From the itemset of size 1 it was noticed that approximately 48.5% of the objects in the dataset are from North America, highlighting the prevalence of this region.

---

[1]see subsection 1.1.3 for the list of the features

The size-2 itemset reveals that 1/4 of the objects is both from North America and has a very low runtime. This pattern becomes more specific in the top size-3 itemset, where almost 10% of the data corresponds to TV episodes with both characteristics.

Size	Support	Itemsets
1	0.485	(is_from_NA)
2	0.204	(VeryLowRTM, is_from_NA)
3	0.094	(tvEpisode, VeryLowRTM, is_from_NA)

Table 5.2: Top itemsets of sizes 1, 2, and 3

5.2 Extraction of rules

After extracting the frequent patterns, association rules were generated. To find a value of *confidence* that balances the number of rules and their strength (measured with *lift*), the plot in Figure5.1b was analysed. A `min_confidence` of 0.55 was selected, guaranteeing an average *lift* of 2.3 and a significant number of rules, i.e. 18. The top 10 rules extracted (ranked by lift) are:

Rule	Antecedents	Consequents	Ant. Sup.	Cons. Sup.	Sup.	Conf.	Lift
0	(short)	(VeryLowC, VeryLowRTM)	0.160	0.128	0.093	0.583	4.568
1	(VeryLowC, VeryLowRTM)	(short)	0.128	0.160	0.093	0.731	4.568
2	(HighRTM)	(movie)	0.177	0.319	0.154	0.866	2.719
3	(is_from_NA, short)	(VeryLowRTM)	0.082	0.375	0.082	0.995	2.654
4	(VeryLowC, short)	(VeryLowRTM)	0.094	0.375	0.093	0.994	2.650
5	(short)	(VeryLowRTM)	0.160	0.375	0.159	0.993	2.648
6	(VeryLowRTM, short)	(VeryLowC)	0.159	0.234	0.093	0.588	2.513
7	(short)	(VeryLowC)	0.160	0.234	0.094	0.587	2.511
8	(is_from_NA, LowRTM)	(tvEpisode)	0.121	0.303	0.090	0.742	2.447
9	(MediumRTM)	(movie)	0.229	0.319	0.165	0.720	2.260

Table 5.3: Top 10 rules extracted with *Apriori* (ranked by lift)

5.3 Exploiting rules for target prediction

One way to exploit the previously extracted rules is for target prediction. Firstly, rules with `VeryLowC` as target (rows 6 and 7 in the Table 5.3), show that short contents with very low runtime are highly likely to be associated with very low credits, probably reflecting the involvement of a limited production or cast. Both rules show a strong association, with a *lift* greater than 2.5, with the more specific one (VeryLowRTM and short) offering a better potential for targeted prediction.

The target `is_from_NA` was then analysed to find the antecedents (as shown in Table 5.4), that increase the likelihood of an object of being from North America. Even though the *lift* value of those rules is below average, these rules were still considered, due to their meaningful interpretability. The analysis suggests that North American origin is associated with TV episodes, shorter durations, and high ratings or numerous production credits.

Rule	Antecedents	Consequents	Ant. Sup.	Cons. Sup.	Sup.	Conf.	Lift
12	(HighR, tvEpisode)	(is_from_NA)	0.144	0.485	0.092	0.639	1.317
13	(HighC)	(is_from_NA)	0.152	0.485	0.156	0.627	1.292
14	(tvEpisode, LowRTM)	(is_from_NA)	0.144	0.485	0.090	0.624	1.285
15	(tvEpisode)	(is_from_NA)	0.303	0.485	0.186	0.612	1.261
16	(tvEpisode, VeryLowRTM)	(is_from_NA)	0.153	0.485	0.093	0.611	1.259
17	(LowRTM)	(is_from_NA)	0.219	0.485	0.121	0.553	1.139

Table 5.4: `is_from_NA` as target